

Assignment 7 | Neha Chede

1.

Input Tables:

```
create table PC (p varchar(30), c varchar(30));
insert into PC (p, c) values
    ('Jay', 'Claire'),
    ('Jay', 'Mitchell'),
    ('Mitchell', 'Lily'),
    ('Claire', 'Haley'),
    ('Claire', 'Alex'),
    ('Claire', 'Luke');
```

```
create table Male (p varchar(30));
insert into Male (p) values
    ('Jay'),
    ('Mitchell'),
    ('Luke');
```

```
create table Female (p varchar(30));
insert into Female (p) values
    ('Claire'),
    ('Lily'),
    ('Haley'),
    ('Alex');
```

Query:

```
with recursive ancestors(ance, id) as (
    select p as ance, c as id
    from PC
    union
    select a.ance, PC.c
    from ancestors a
    join PC
    on a.id = PC.p
),
descendants(desce, id) as (
    select c as desce, p as id
    from PC
    union
    select a.desce, PC.p
    from descendants a
    join PC
    on a.id = PC.c
)
select a.ance as x, m.p as y, d.desce as z
from ancestors a, descendants d, Male m, Female f
where a.id = m.p and m.p = d.id and d.desce = f.p;
```

Output:

x	y	z
Jay	Mitchell	Lily


2.

$$|A - B| = 0$$

$$\neg \exists x (x \in A - B)$$

A: Number of people known by Person, p1

B: Number of people known by Person, p2



Query:

```
create or replace function k_cnt(pid int)
returns int as
$$
    select count(k.pid2)
    from Knows k
    where k.pid1 = k_cnt.pid
    group by k.pid1;
$$ language sql;

select distinct k1.pid1 as p1, k2.pid1 as p2
from Knows k1, Knows k2
where k1.pid1 <> k2.pid1 and not exists(
    select *
    from k_cnt(k1.pid1)
    except
    select *
    from k_cnt(k2.pid1)
);
```

3.

```
select distinct p1.pid
from personHasSkills p1
where not exists(
    select 1
    from personHasSkills p2
    where cardinality(p2.skills) > cardinality(p1.skills)
      and p1.pid<>p2.pid
);
```

Output:

	pid integer
1	1011

4.

Considering constants, a = 5 and c = 14,

```
create table r(a integer, b integer);
create table s(b integer, c integer);
```

```
create table v_cnt (cnt integer);
insert into v_cnt values (0);
```

```
create or replace function upd_cnt()
returns trigger as
$$
begin
    if tg_op = 'INSERT' then
        if (new.a!=5 and new.b in (select b from s where c!=14)) then
            update v_cnt
            set cnt = cnt + 1;
        end if;
    elsif tg_op = 'DELETE' then
        if (old.a!=5 and old.b in (select b from s where c!=14)) then
            update v_cnt
            set cnt = cnt - 1;
        end if;
    end if;
    return new;
end;
$$ language plpgsql;
```

```
create or replace trigger insert_r after insert on r
for each row
execute function upd_cnt();
```

```

create or replace trigger delete_r after delete on r
for each row
execute function upd_cnt();

insert into r values (1,2), (3,4), (5,6), (7,8), (9,10);
insert into s values (2,11), (4,12), (6,13), (8,14), (10,15);

create or replace view ques as
    SELECT r.a, s.c
    FROM R r, S s
    WHERE r.a != 5 AND r.b = s.b AND s.c != 14;

select * from v_cnt;

```

Output:

```

"cnt"
3

```

```

insert into r values (2,20);
select * from v_cnt;

```

Output:

```

"cnt"
4

```

5. With buffers, $B(R)$, $B(S)$, $B(T)$ and block size, M , if we use the block nested-loop algorithm to implement natural join operations, to evaluate the relational algebra expression $(R \bowtie S) \bowtie T$, the time complexity is:

$$B(R \bowtie S) = B(R) + \frac{B(R) \times B(S)}{M}$$

where $B(R \bowtie S)$ is the number of blocks to store $(R \bowtie S)$.

$$B((R \bowtie S) \bowtie T) = B(R \bowtie S) + \frac{B(R \bowtie S) \times B(T)}{M}$$

Given the assumption, $B(R \bowtie S) \leq M^2$, the overall time complexity depends on the number of block transfers required for each join operation.

6. (a) Given, $r = 300,000$ records, $B = 4,096$ bytes, length of record = 100 bytes

$$\text{Number of Block Accesses} = \lceil \log_2(N + 1) \rceil = \left\lceil \log_2\left(\frac{300000 \times 100}{4096} + 1\right) \right\rceil = \lceil \log_2(7325) \rceil = \lceil 12.84 \rceil = 13$$

6. (b) Given, $V = 9$ bytes, $P = 6$ bytes

$$\text{Records per block} = \left\lfloor \frac{4096}{100} \right\rfloor = 40$$

$$\text{Entries per block, } b = \left\lfloor \frac{4096}{9+6} \right\rfloor = \lfloor 273.06 \rfloor = 273 \text{ indexes}$$


$$\text{Number of blocks in the index file} = \left\lceil \frac{300000/40}{273} \right\rceil = \lceil 27.47 \rceil = 28$$

With binary search as a primary index is constructed, number of block access required = $\lceil \log_2 28 \rceil = \lceil 4.80 \rceil = 5$

7. (a) $R1(x); R2(y); R1(z); R2(x); R1(y)$

The schedule is conflict serializable as there are all read operations involved in the transaction and there are no cycles in the precedence graph. A conflict equivalent serial schedule would be: $R1(x); R1(z); R2(y); R2(x); R1(y)$.

7. (b) $R1(x); W2(y); R1(z); R3(z); W2(x); R1(y)$



There is a cycle in the precedence graph for the given schedule [Transactions (1, 2, 3, 1)], thus, this schedule is not conflict serializable.

8. (a) The relation:

Patients (Patient_ID, Name, DOB)

Doctors (Doctor_ID, name, specialty)

Nurses (Nurse_ID, name, department)

Relationship tables: pat_doc (Patient_ID, Doctor_ID) pat_nur (Patient_ID, Nurse_ID)

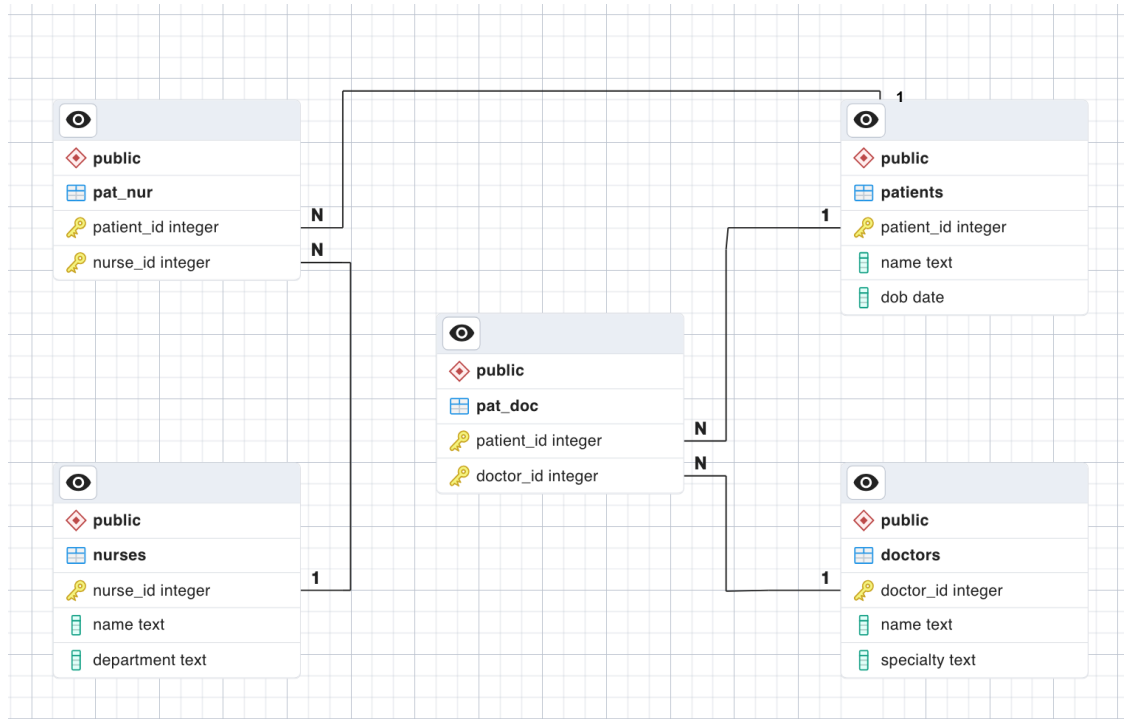


Fig. Entity Relationship Diagram for the given schema using pgAdmin

8. (b)

```

create table Patients(
    Patient_ID integer primary key,
    name text,
    DOB date
);

create table Doctors(
    Doctor_ID integer primary key,
    name text,
    specialty text
);

create table Nurses(
    Nurse_ID integer primary key,
    name text,
    department text
);

create table pat_doc(
    Patient_ID integer references Patients(Patient_ID),
    Doctor_ID integer references Doctors(Doctor_ID),
    primary key(Patient_ID, Doctor_ID)
);

create table pat_nur(

```

```
Patient_ID integer references Patients(Patient_ID),  
Nurse_ID integer references Nurses(Nurse_ID),  
primary key(Patient_ID, Nurse_ID)  
);
```