

The Object-Relational Database Model- an Entry in noSQL database models

Dirk Van Gucht¹

¹Indiana University

March 9, 2019

Motivation

- In the relational database model, each relation has a schema of attributes with **atomic domains** such as booleans, numbers, text strings etc
- In the **object-relational model**, and in various **noSQL data models**, attributes in the schema the database may also have domains that consists of **complex objects** such as arrays, sets, bags, objects of composite types, relations, **JSON** objects, **XML** documents etc
- Therefore we need mechanisms to
 - 1 define relations/databases with attributes of complex-object types; and
 - 2 search and manipulate such relations/databases

Set and bag types as arrays

- In general, complex-object types can be recursively defined in terms of atomic types, composite types, array types, etc
- The main focus of this lecture will be on **array types**
- In particular, we will show how array types can be used to model **bag** and **set** types
- We will then show how operations on arrays allow us to model operations on bags and sets
- **PostgreSQL** is an excellent system to consider the issues since it is an object-relational database system
- Many of the concepts discussed here can also be found in the **noSQL MongoDB** system as well as in the **MapReduce** framework and its derivatives

Arrays and the array constructor operation

- In SQL,

`{7, 4, 4, 3, 2}'::int[]`

denotes the array [7, 4, 3, 3, 2] of type `int[]`

- Its first, third, and fifth component values are obtained as follows:

Array component	Value
<code>{7, 4, 4, 3, 2}'::int[]</code> [1]	7
<code>{7, 4, 4, 3, 2}'::int[]</code> [3]	4
<code>{7, 4, 4, 3, 2}'::int[]</code> [5]	2

- In SQL,

`{“C”,“John”,“Anna”,“12”}'::text[]`

denotes the array ['C', 'John', 'Anna', '12'] of type `text[]`

- Elements of an array must all be of the same type

Arrays and the array constructor operation (alternative syntax)

- In SQL, the following all denote the same array of integers

[7, 4, 3, 3, 2]

{7, 4, 4, 3, 2}::int[]

ARRAY[7,4,4,3,2]::int[]

ARRAY[7,4,4,3,2]

- The third component value is obtained as follows

Array component	Value
({7, 4, 4, 3, 2}::int[])[3]	4
(ARRAY[7,4,4,3,2]::int[])[3]	4
(ARRAY[7,4,4,3,2])[3]	4

Modeling bags and sets with arrays

- The array

`ARRAY[7,4,4,3,2]`

represent (models) the bag

`{2, 3, 4, 4, 7}`

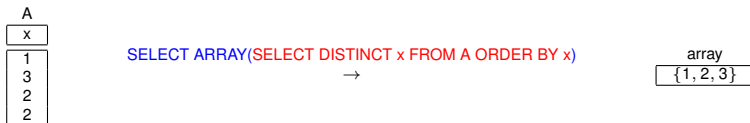
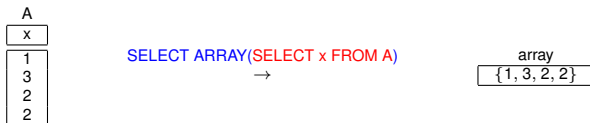
and the set

`{2, 3, 4, 7}`

- Recall that an array orders its elements but a bag or a set does not
- The arrays `ARRAY[7,4,4,3,2]` and `ARRAY[2,4,3,7,4]` are **different** but they both represent the **same** bag and the **same** set
- The empty array `{}` or `ARRAY[]` models the empty set `{}` (i.e., \emptyset)

ARRAY construction from a unary SQL query

- The ARRAY constructor operation can be applied to any SQL query that returns a **unary** relation
- It constructs an array of the elements of that relation



ARRAY construction from a SQL query with ROW construction

- The ARRAY constructor operation can be applied to any SQL query
- But, the tuples returned by the query must be **packed** to be of a composite (row) type by the **ROW** constructor operation

A	
x	y
1	a
1	b
2	a

SELECT ARRAY(SELECT ROW(x,y) FROM A)

→
SELECT ARRAY(SELECT (x,y) FROM A)

array
{(1, a), (1, b), (2, a)}

Example: A documents relation

- We may wish to maintain a database of documents and the words they contain
- We can define a relation with attributes of atomic domain to store such documents

`CREATE TABLE documentWord (doc text, word text);`

- A pair (d, w) in `documentWord` specifies that document d contains the word w

Example: A documents relation

- The documentWord relation may look as follow:

documentWord

doc	word
d1	A
d1	B
d1	C
d2	B
d2	C
d2	D
d3	A
d3	E
d4	B
d4	B
d4	A
d4	D
d5	E
d5	F
d6	A
d6	D
d6	G
d7	C
d7	B
d7	A
d8	B
d8	A

Example: The documents relation as a complex-objects relation

- We could consider a more **natural representation** of this data by having a relation **documents** of pairs (doc, words) where we pair each document with its set (bag) of words

documents

L

doc	words
d1	{A, B, C}
d2	{B, C, D}
d3	{A, E}
d4	{B, B, A, D}
d5	{E, F}
d6	{A, D, G}
d7	{C, B, A}
d8	{B, A}

- Such a relation is called a **complex-objects relation**
- We will see how the **ARRAY** type can be used to model such complex-objects relations
- We will then discuss how such relations can be manipulated and queried

The ARRAY type

- SQL has the array type
 - 1 for example, the array type `text[]` declares an array of text;
 - 2 `int[]` declares an array of int;
- SQL permits the use of these types in the definition of complex-object relations. For the `documents` relation, we can use the declaration

```
CREATE TABLE documents (doc text, words text[]);
```

- Such a table can be populated using insert statements such as

```
...  
INSERT INTO documents VALUES ('d6', {'A','D','G'});  
...
```

Querying the documents relation

- Next consider the query

`SELECT d.doc, d.words FROM documents d`

- This query returns the contents of the documents relation
- The result would be

doc	words
d1	{A, B, C}
d2	{B, C, D}
d3	{A, E}
d4	{B, B, A, D}
d5	{E, F}
d6	{A, D, G}
d7	{C, B, A}
d8	{B, A}

Set and Bags as Unordered Array

- We will use arrays to represent sets (or bags).
- We must therefore restrict the predicates and operations we define on arrays to be **independent** of the order in which the elements appear in the arrays
- The following are such predicates and operations

$a \in A$

a is a **member** (element) of set A

$a \notin A$

a is **not a member** (element) of set A

$A \cap B \neq \emptyset$

set A and set B **overlap**

$A \subseteq B$

set A is a **subset** of set B

$A \supseteq B$

set A is a **superset** of set B

$A = \emptyset$

set A is **empty**

$|A|$

denotes the **cardinality** (size) of set A

$A \cup B, A \cap B, A - B$

union, intersection, difference

of sets A and B

Checking for **Set Membership** ($a \in A$)

- In SQL this can be done using the = **SOME** predicate
- “Find the documents that **contain** the word ‘D’ ”

```
SELECT  d.doc, d.words  
FROM    documents d  
WHERE   'D' = SOME(d.words)
```

doc	words
d2	{ B, C, D }
d4	{ B, B, A, D }
d6	{ A, D, G }

Checking for **Set Non-Membership** ($a \notin A$)

- “Find the documents that do **not contain** the word ‘D’ ”
- For this we can use the **<> ALL** predicate.

```
SELECT  d.doc, d.words
FROM    documents d
WHERE   ‘D’ <> ALL(d.words)
```

doc	words
d1	{A, B, C}
d3	{A, E}
d5	{E, F}
d7	{C, B, A}
d8	{B, A}

The `isIn` set-membership function

- For convenience, we define a **polymorphic function** `isIn` for the set-membership predicate:

```
CREATE FUNCTION isIn (x anyelement, A anyarray)
  RETURNS boolean AS
  $$
    SELECT x = SOME(A);
  $$ LANGUAGE SQL;
```

- We can now write the query “Find the documents that contain the word ‘A’ but not the word ‘D’ ” as follows:

```
SELECT  d.doc, d.words
FROM    documents d
WHERE   isIn(‘A’,d.words) and not(isIn(‘D’,d.words))
```

Checking for **Overlap of Sets** ($A \cap B \neq \emptyset$)

- We may wish to check if sets **overlap**, i.e., if they have a **non-empty intersection**
- This can be done using the **&&** predicate.
- “Find the documents whose sets of words overlap with the set of words $\{B, C\}$.”

```
SELECT  d.doc, d.words
FROM    documents d
WHERE   d.words && '{"B","C"}'
```

doc	words
d1	{A, B, C}
d2	{B, C, D}
d4	{B, B, A, D}
d7	{C, B, A}
d8	{B, A}

Checking for **Disjoint (Non-overlapping) Sets** ($A \cap B = \emptyset$)

- “Find the pairs of documents that do not have words in common.”

```
SELECT  d1.doc AS doc1, d2.doc AS doc2,  
        d1.words AS words1, d2.words AS words2  
FROM    documents d1, documents d2  
WHERE   NOT( d1.words && d2.words )
```

doc1	doc2	words1	words2
d1	d5	{A, B, C}	{E, F}
d2	d3	{B, C, D}	{A, E}
d2	d5	{B, C, D}	{E, F}
d3	d2	{A, E}	{B, C, D}
d4	d5	{B, B, A, D}	{E, F}
d5	d1	{E, F}	{A, B, C}
d5	d2	{E, F}	{B, C, D}
d5	d4	{E, F}	{B, B, A, D}
d5	d6	{E, F}	{A, D, G}
d5	d7	{E, F}	{C, B, A}
d5	d8	{E, F}	{B, A}
d6	d5	{A, D, G}	{E, F}
d7	d5	{C, B, A}	{E, F}
d8	d5	{B, A}	{E, F}

Checking for **Set Containment (subset)** ($A \subseteq B$)

- We may wish to check if a set is a **subset** of another set
- This can be done using the **<@** set-containment predicate
- “Find the documents that contain the words ‘A’ and ‘B’ ”

```
SELECT  d.doc, d.words
FROM    documents d
WHERE   {'A', 'B'} <@ d.words
```

doc	words1
d1	{A, B, C}
d4	{B, B, A, D}
d7	{C, B, A}
d8	{B, A}

Checking for Set Containment (subset)

- “Find the pairs of different documents d1, d2 such that all words in d1 also occur as words in d2.”

```
SELECT  d1.doc AS doc1, d2.doc AS doc2,  
        d1.words AS words1, d2.words AS words2  
FROM    documents d1, documents d2  
WHERE   d1.words <@ d2.words AND  
        d1.doc <> d2.doc
```

doc1	doc2	words1	words2
d1	d7	{A, B, C}	{C, B, A}
d7	d1	{C, B, A}	{A, B, C}
d8	d1	{B, A}	{A, B, C}
d8	d4	{B, A}	{B, B, A, D}
d8	d7	{B, A}	{C, B, A}

Checking for Set Equality ($A = B$)

- We may wish to check if two sets are **equal**
- This can again be done using the **<@** set-containment predicate
- "Find the pairs of different documents d1, d2 that have the same words."

```
SELECT  d1.doc AS doc1, d2.doc AS doc2,  
        d1.words AS words1, d2.words AS words2  
FROM    documents d1, documents d2  
WHERE   d1.words <@ d2.words AND  
        d2.words <@ d1.words AND  
        d1.doc <> d2.doc
```

doc1	doc2	words1	words2
d1	d7	{A, B, C}	{C, B, A}
d7	d1	{C, B, A}	{A, B, C}

Caveat: Do not use ARRAY equality = to test set-equality

- Consider the ARRAY equality predicate '='
- This predicate checks if two arrays are the same, i.e., they are equal **component by component**
- So '=' is an **order-dependent** predicate and should therefore not be used in our context of set predicates and operations

```
SELECT  d1.doc AS doc1, d2.doc AS doc2,  
        d1.words AS words1, d2.words AS words2  
FROM    documents d1, documents d2  
WHERE   d1.words = d2.words AND  
        d1.doc <> d2.doc
```

will return the empty set

Checking for Set Emptiness ($A = \emptyset$)

- “Find the documents that contain no words.”

```
SELECT  d.doc, d.words
FROM    documents d
WHERE   d.words <@ '{}'
```

- Recall that '{}' represents the empty set

Application: Set joins

- Recall queries of the form: “Find all pairs of documents (d_1, d_2) such that **some** | **not all** | **not only** | **no** | **all** | **only** words of d_1 are in d_2 .”
- These set-joins can be captured using the **overlap** and **containment** predicates
- To do so, we can define polymorphic user-defined functions that stand for these set-join predicates
- We will illustrate this for the **some** (i.e., **at least one**) and **all** set joins. The other set joins can be specified in a similar fashion

Application: Set joins

- **SOME** (at least one) set join

```
CREATE OR REPLACE FUNCTION atLeastOne (A anyarray, B anyarray)
RETURNS boolean AS
$$
    SELECT A && B;
$$ LANGUAGE SQL;
```

- **ALL** set join (better called SUBSET join)
"Is **each** element in *A* an element of *B*?"

```
CREATE OR REPLACE FUNCTION Each (A anyarray, B anyarray)
RETURNS boolean AS
$$
    SELECT A <@ B;
$$ LANGUAGE SQL;
```

Application: Set joins

We can then write queries with set joins as follows:

- “Find all pairs of documents (d_1, d_2) such that **some** words of d_1 are in d_2 .”

```
SELECT d1.doc, d2.doc  
FROM documents d1, documents d2  
WHERE atLeastOne(d1.words,d2.words)
```

- “Find all pairs of documents (d_1, d_2) such that **all** words of d_1 are in d_2 .”

Alternatively, “Find all pairs of documents (d_1, d_2) such that d_1 only contains words that are in d_2 .”

```
SELECT d1.doc, d2.doc  
FROM documents d1, documents d2  
WHERE Each(d1.words,d2.words)
```

Determining Set Size (Cardinality) ($|A|$)

- We may wish to determine the size (cardinality) of sets
- This can be done using the ARRAY **cardinality** function
- “Find the number of words in each document.”

```
SELECT  d.doc, cardinality(d.words) AS number_of_words  
FROM    documents d
```

doc	number_of_words
d1	3
d2	3
d3	2
d4	4
d5	2
d6	3
d7	3
d8	2

Example: Queries using set cardinality

- “Find the documents with fewer than 10 words”

```
SELECT  d.doc  
FROM    documents d  
WHERE   cardinality(d.words) < 10
```

The UNNEST operator

- It is possible to **coerce** an array into a (unary) relation that contains the elements of the array
- This is done using the **UNNEST** operator

`SELECT UNNEST(ARRAY[2,1,3,4,4])` →

unnest
2
1
3
4
4

- It is possible to provide an attribute name for the elements

`SELECT UNNEST(ARRAY[2,1,3,4,4]) AS A` →

A
2
1
3
4
4

Restructuring: the UNNEST operator

- It is possible to **restructure** a complex-object relation by using the **UNNEST** restructuring operator
- "Starting from the **documents** relation, create a relation of (doc, word) pairs."

```
SELECT  d.doc, UNNEST(d.words) AS word  
FROM    documents d
```

→

doc	word
d1	A
d1	B
d1	C
d2	B
d2	C
d2	D
d3	A
d3	E
d4	B
d4	B
d4	A
d4	D
d5	E
d5	F
d6	A
d6	D
d6	G
d7	C
d7	B
d7	A
d8	B

Set operations: setUnion, setIntersection, and setDifference

- Using UNNEST and ARRAY construction it is also possible to define **setUnion**, **Intersection**, and **Difference** on sets represented as arrays
- We do this with polymorphic functions.
- Here we will show how to do this for **setUnion**

```
CREATE FUNCTION setUnion (A anyarray, B anyarray) RETURNS anyarray AS
$$
SELECT ARRAY( SELECT * FROM UNNEST(A)
               UNION
               SELECT * FROM UNNEST(B));
$$ LANGUAGE SQL;
```

```
SELECT setUnion( '{1, 2, 3}'::int[], '{2, 3, 3, 5}'::int[] );
```

setUnion
{1, 2, 3, 5}

```
SELECT setUnion( '{" A" , " B" }'::text[], '{" A" , " C" }'::text[] );
```

setUnion
{A, B, C}

Restructuring: GROUPING (nesting)

- Reconsider the documentWord relation
- “Restructure this relation by grouping the words of each document into a set (bag)”

documentWord

doc	word
d1	A
d1	B
d1	C
d2	B
d2	C
d2	D
d3	A
d3	E
d4	B
d4	B
d4	A
d4	D
d5	E
d5	F
d6	A
d6	D
d6	G
d7	C
d7	B
d7	A
d8	B

group words by doc



doc	words
d1	{A, B, C}
d2	{B, C, D}
d3	{A, E}
d4	{B, B, A, D}
d5	{E, F}
d6	{A, D, G}
d7	{C, B, A}
d8	{B, A}

Restructuring: GROUPING (nesting)

- This can be done using the **ARRAY** constructor operation

```
SELECT  DISTINCT d.doc,  
        ARRAY(SELECT d1.word  
                FROM   documentWord d1  
                WHERE  d1.doc = d.doc) AS words  
FROM    documentWord d;
```

- Notice how the parameter **d** is used inside the **ARRAY** constructor to group together the words associated with the document **d**
- The **DISTINCT** operation is essential
- This query runs in $O(|documentWord|^2)$.

Restructuring: GROUPING (nesting) using the array_agg function

- The same restructuring can also be done using the `array_agg` aggregate function

```
SELECT d.doc, array_agg(d.word)
FROM   documentWord d
GROUP BY (d.doc)
```

- The `GROUP BY(d.doc)` operation partitions the `documentWord` by `doc` values
- For each cell in this partition, the `array_agg` function aggregates in an array the words that are in that cell
- This query runs in $O(|documentWord|)$
- So much faster than the other restructuring query

Repeated restructuring (Different views of same data)

- Starting from the **documents** relation, we may want to create a complex-object relation **words** which keeps for each word the set of documents that contain that word
- In other words, we want to do the following restructuring

doc	words
d1	{A, B, C}
d2	{B, C, D}
d3	{A, E}
d4	{B, B, A, D}
d5	{E, F}
d6	{A, D, G}
d7	{C, B, A}
d8	{B, A}

restructure:
Step 1: unnest on words;
Step 2: group docs by word
→

word	docs
A	{d1, d3, d4, d6, d7, d8}
B	{d1, d2, d4, d7, d8}
C	{d1, d2, d7}
D	{d2, d4, d6}
E	{d3, d5}
F	{d5}
G	{d6}

- This can be accomplished by unnesting the **documents** relation on words and then grouping the doc values by word

Repeated restructuring

```
WITH docWord AS (SELECT d.doc AS doc,  
                        UNNEST(d.words) AS word  
                  FROM documents d)  
SELECT p.word AS word, array_agg(p.doc) AS docs  
FROM   docWord p  
GROUP BY (p.word)
```

Or, as one query

```
SELECT      word, array_agg(doc) AS docs  
FROM        (SELECT doc, UNNEST(words) AS word  
              FROM   documents d) p  
GROUP BY    (word)
```

Application: The word-count problem

“Determine the word-count, i.e., frequency of occurrence, of each word in the set of documents”

```
SELECT      word, cardinality(array_agg(doc)) AS wordCount
FROM        (SELECT doc, UNNEST(words) AS word
              FROM    documents d) p
GROUP BY    (word)
```

doc	words
d1	{A, B, C}
d2	{B, C, D}
d3	{A, E}
d4	{B, B, A, D}
d5	{E, F}
d6	{A, D, G}
d7	{C, B, A}
d8	{B, A}

→

word	wordCount
F	1
G	1
E	2
C	3
D	3
A	6
B	6

Application: The most frequent words

“Find the words that occur most frequently in the set of documents.”

```
WITH E AS (  
  SELECT word, cardinality(array_agg (doc)) AS wordCount  
  FROM (SELECT doc, UNNEST(words) AS word  
        FROM documents d) p  
  GROUP BY (word))  
  
SELECT word  
FROM E  
WHERE wordCount = (SELECT MAX(wordCount) FROM E)
```

Double nesting

- Consider the following **Enroll(sid,cno,grade)** relation

sid	cno	grade
1001	2001	A
1001	2002	A
1001	2003	B
1002	2001	B
1002	2003	A
1003	2004	A
1003	2005	B
1004	2002	A
1004	2004	A
1005	2001	B
1005	2003	A

- From this we want to create a complex-object relation which stores for each student, his or her courses, internally grouped by grades obtained in these courses
- This requires double nesting

Double nesting

- We begin by grouping on (sid,grade)

```
SELECT e.sid, e.grade, array_agg(e.cno) AS courses
FROM enroll e
GROUP BY (e.sid, e.grade)
```

- This gives the complex-object relation

sid	grade	courses
1001	A	{2001, 2002}
1001	B	{2003}
1002	A	{2003}
1002	B	{2001}
1003	A	{2004}
1003	B	{2005}
1004	A	{2002, 2004}
1005	A	{2003}
1005	B	{2001}

Double nesting

- We then group over the pair of attributes (grade,courses)

```
WITH F AS (SELECT e.sid, e.grade, array_agg(e.cno) AS courses
            FROM enroll e
            GROUP BY (e.sid, e.grade))
```

```
SELECT f.sid, array_agg((f.grade, f.courses)) AS grades
FROM   F f
GROUP BY (f.sid)
```

- Notice the clause array_agg((f.grade,f.course))
- Recall that it is required to make a row (f.grade,f.course) since the array_agg function can only make an array wherein the array values are single values
- I.e., it is not allowed to write array_agg(e.grade,e.course)

Double nesting

sid	cno	grade
1001	2001	A
1001	2002	A
1001	2003	B
1002	2001	B
1002	2003	A
1003	2004	A
1003	2005	B
1004	2002	A
1004	2004	A
1005	2001	B
1005	2003	A

group by (cno)
group by(grade, courses)
→

sid	grades
1001	{ "(A, "{2001, 2002}")", "(B, "{2003}")" }
1002	{ "(A, "{2003}")", "(B, "{2001}")" }
1003	{ "(A, "{2004}")", "(B, "{2005}")" }
1004	{ "(A, "{2002, 2004}")" }
1005	{ "(A, "{2003}")", "(B, "{2001}")" }

- For example, student 1001 obtained two types of grades: 'A' and 'B'
- She received an 'A' in courses 2001 and 2002, and a 'B' in course 2003