

Hashing

Dirk Van Gucht

Based on slides by Hector Garcia-
Molina

Hashing: hash function

- Let K be a domain of key values (K can be very large)
- Let $R = [0, m)$ (usually $m \ll |K|$) be a range of values
- A **hash function** h maps K to R

$$h: K \rightarrow R$$

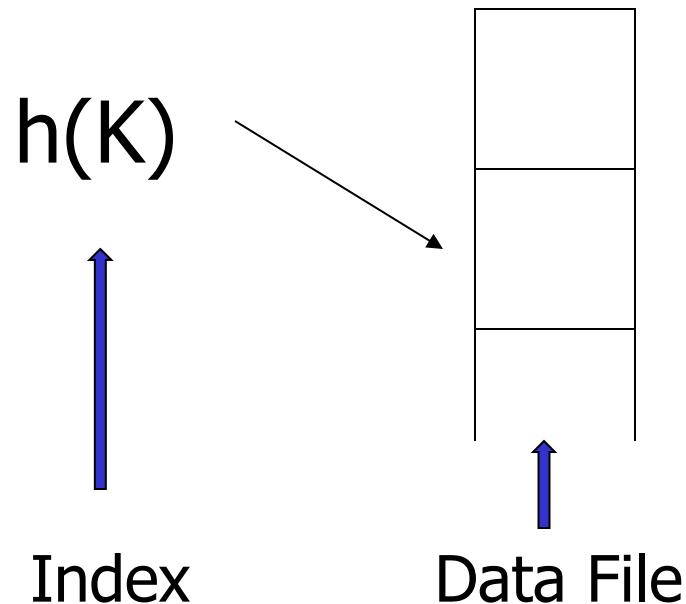
Example: $h(k) = k \bmod m$

- Given a record r with key value k , $h(k)$ provides an address (name) of a **bucket** in which to store r
- A bucket is stored in secondary memory as a block or a list of blocks
- Retrieving a bucket (and the records therein) can be done in $O(n)$ where n is the number of blocks that store the bucket.

Hash function: collision

- Let r_1 and r_2 be two records with key values k_1 and k_2
- We permit that $k_1 = k_2$
- We say that h has a collision for r_1 and r_2 if $h(k_1) = h(k_2)$
Consequence: r_1 and r_2 will be stored in the same bucket
- If K is the domain of a primary key, a collision will store records with different key values in same bucket (not desirable)
- If K is not the domain of a primary key, different records with the same key values will be placed in the same bucket (desirable). Partitioning.
- The latter property is exploited in key-value stores since records with the same key value will be sent to the same reducer

Next: example to illustrate
inserts, overflows, deletes



EXAMPLE 2 records/bucket

(For simplicity, we identify a record with its key value)


INSERT:

$h(a) = 1$

$h(b) = 2$

$h(c) = 1$

$h(d) = 0$



EXAMPLE 2 records/bucket

INSERT:


$h(a) = 1$

$h(b) = 2$

$h(c) = 1$

$h(d) = 0$

$h(e) = 1$



EXAMPLE 2 records/bucket

INSERT:


$h(a) = 1$

$h(b) = 2$

$h(c) = 1$

$h(d) = 0$

$h(e) = 1$




EXAMPLE: deletion

Delete:

e

f




EXAMPLE: deletion

Delete:

e

f

c




EXAMPLE: deletion

Delete:

e

f

c




How do we cope with growth?

- Overflows and reorganizations
Reorganization can be done by enlarging the range R and changing the hash function
Reorganization requires complete rehashing and is linear in the $|Data\ file|$
- Dynamic hashing (extensible hashing)


Extensible hashing: two ideas

(a) Use i of b bits output by hash function



use $i \rightarrow$ grows/shrinks over time....


(b) Use directory of size 2^i



$h(k)$ has b bits, but we will only look at its first i bits

$h(k)[i]$ consists of the first i bits of $h(k)$
these i bits specify the position for k the
directory


Example: $h(k)$ is 4 bits; 2 keys/bucket




local depth of bucket $\leq i$

Insert 1010

Example: $h(k)$ is 4 bits; 2 keys/bucket



Insert 1010



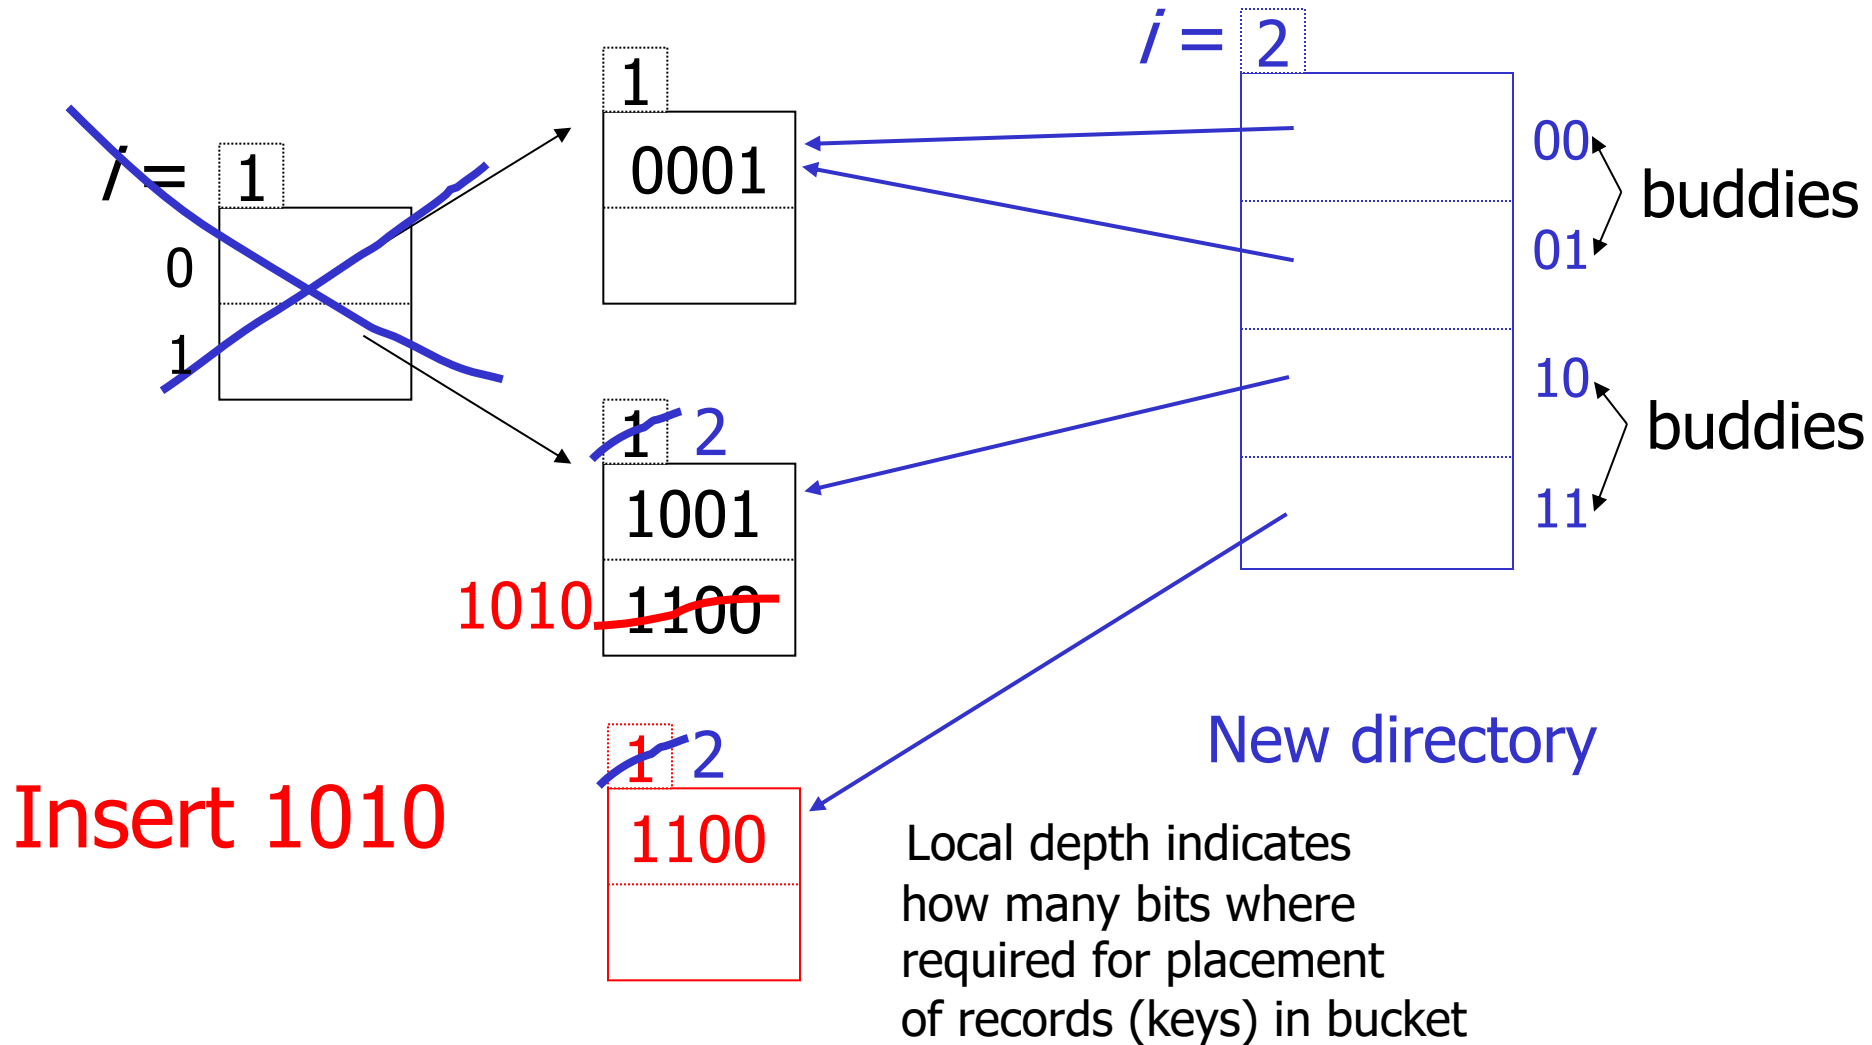
Main idea:
we need 2 bits to
distinguish

1001

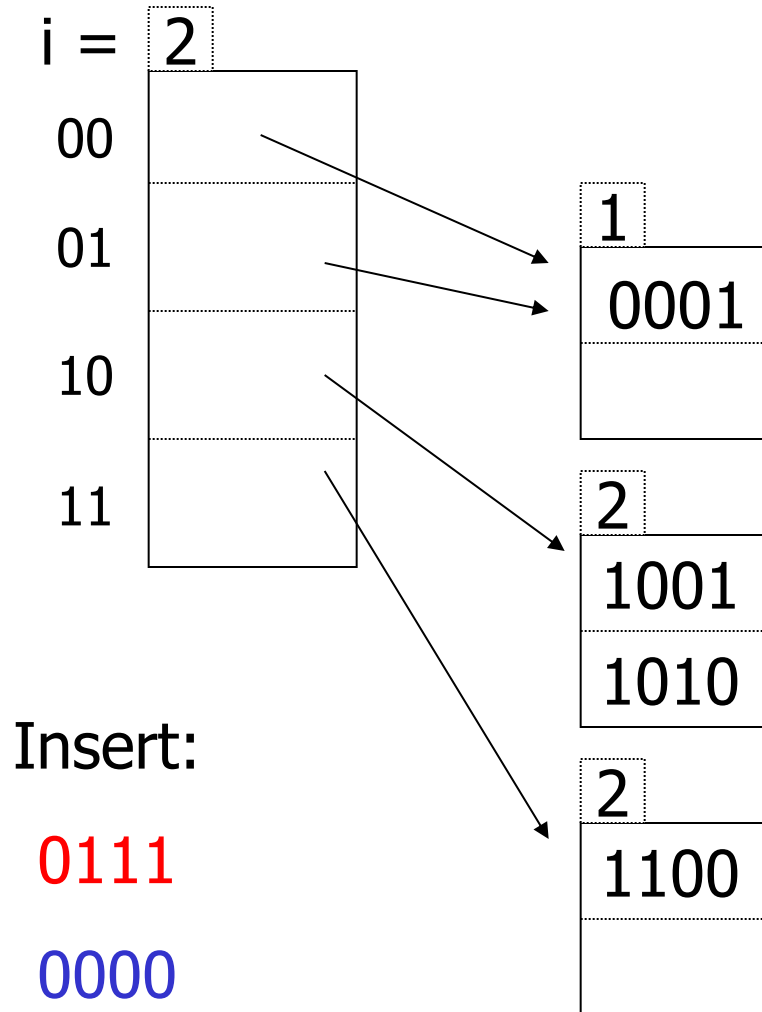
1100

1010

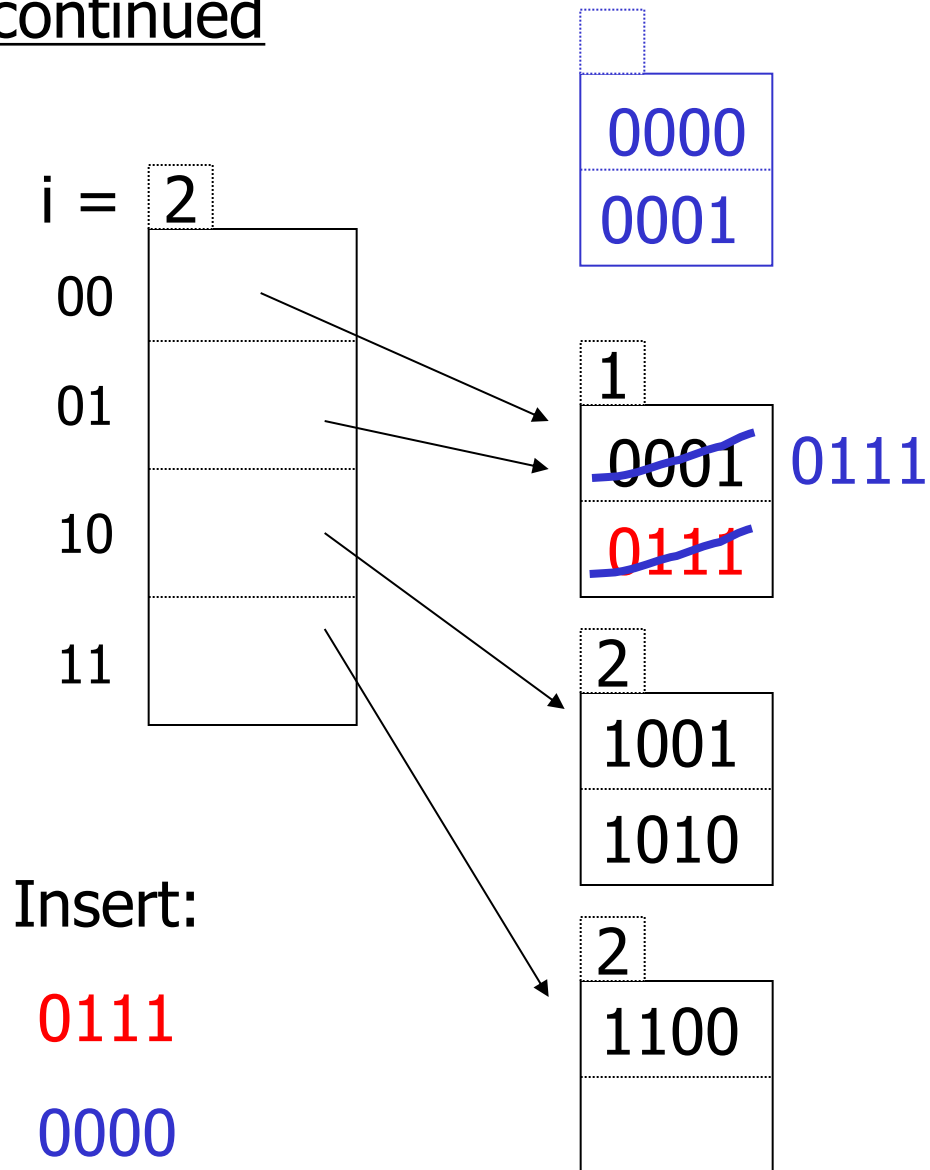
Example: $h(k)$ is 4 bits; 2 keys/bucket



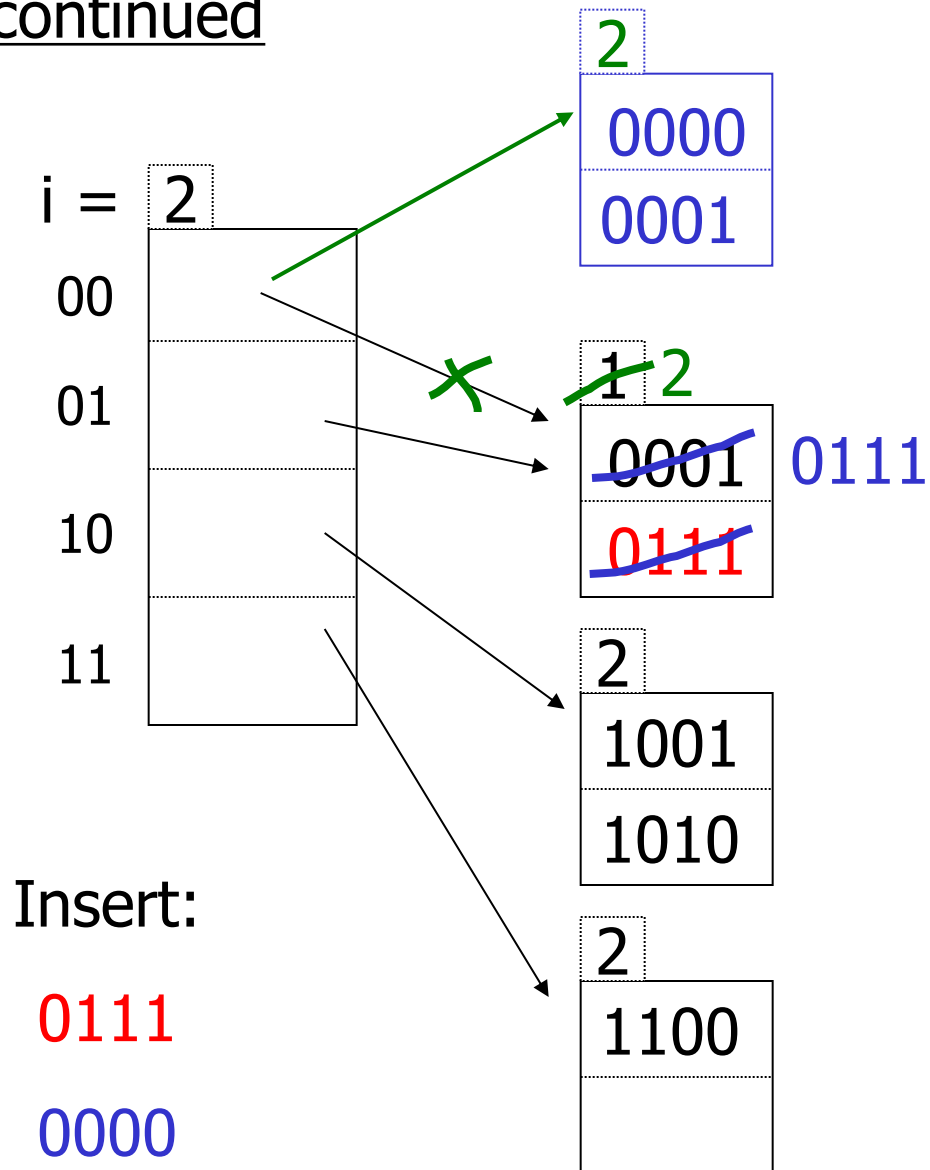
Example continued



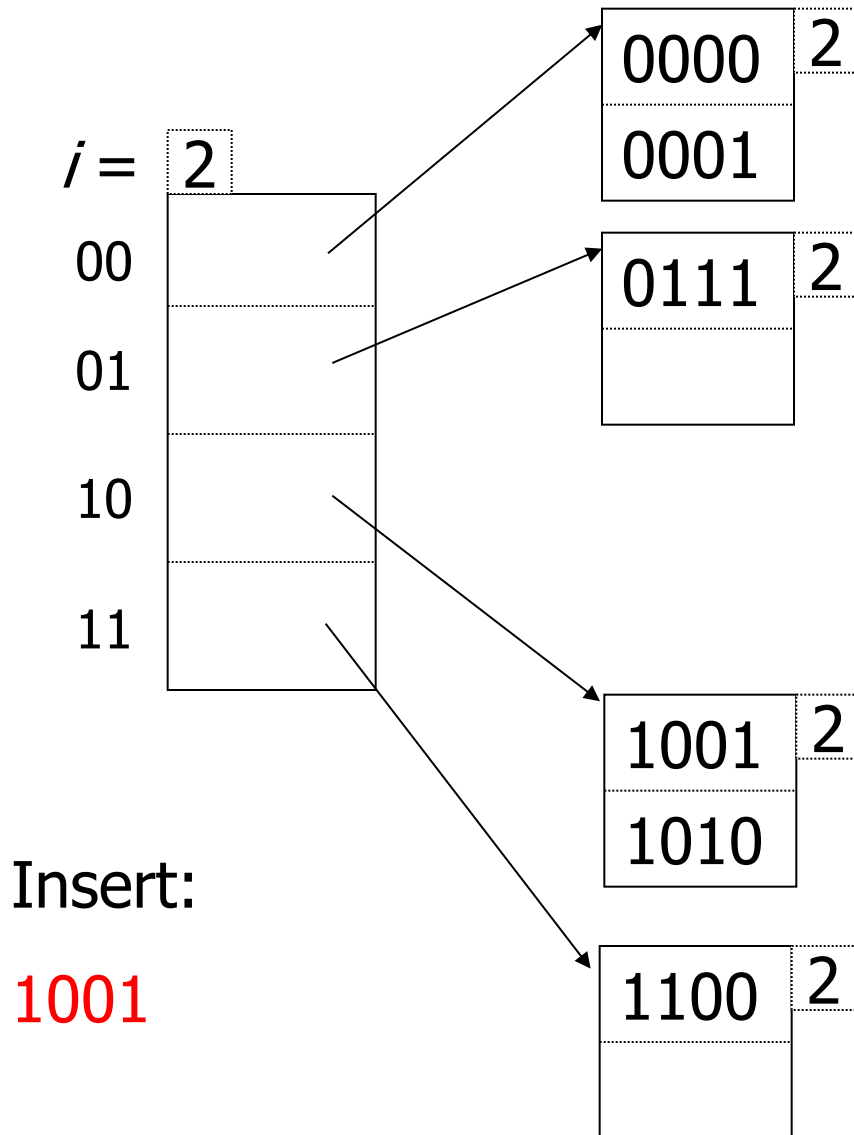
Example continued



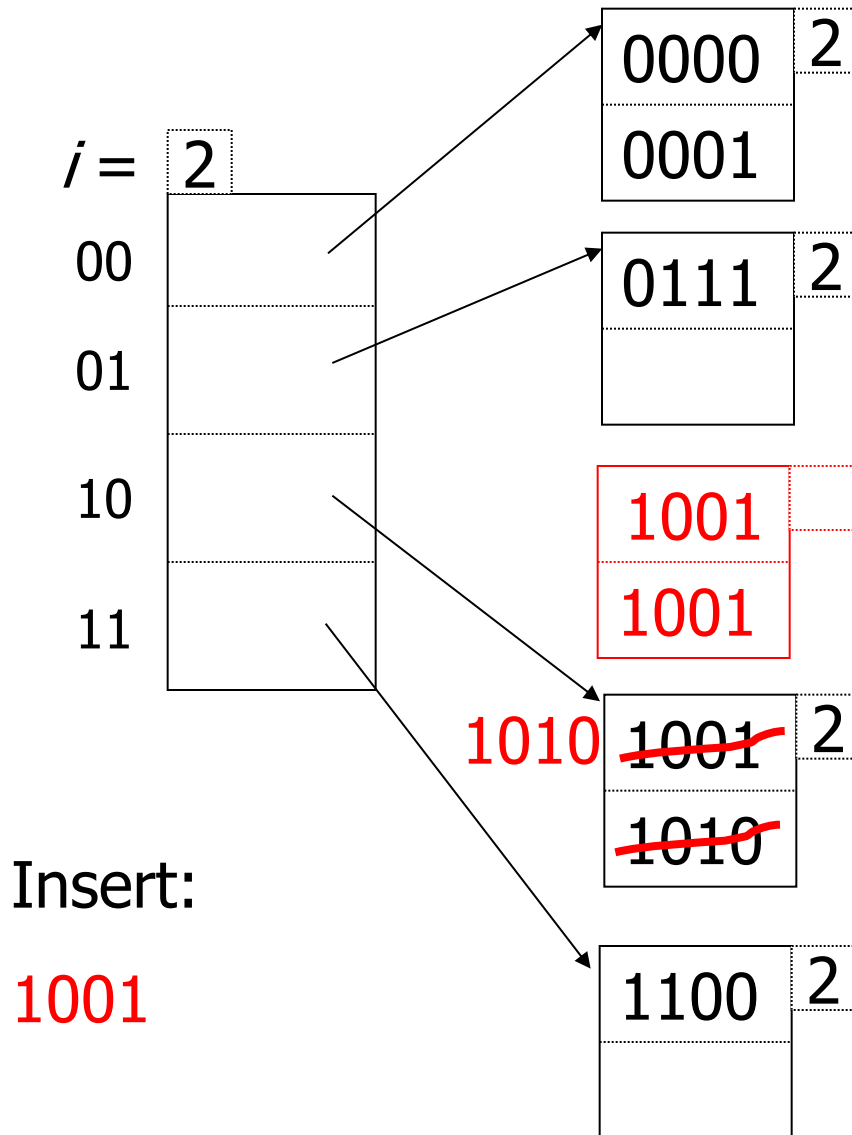
Example continued



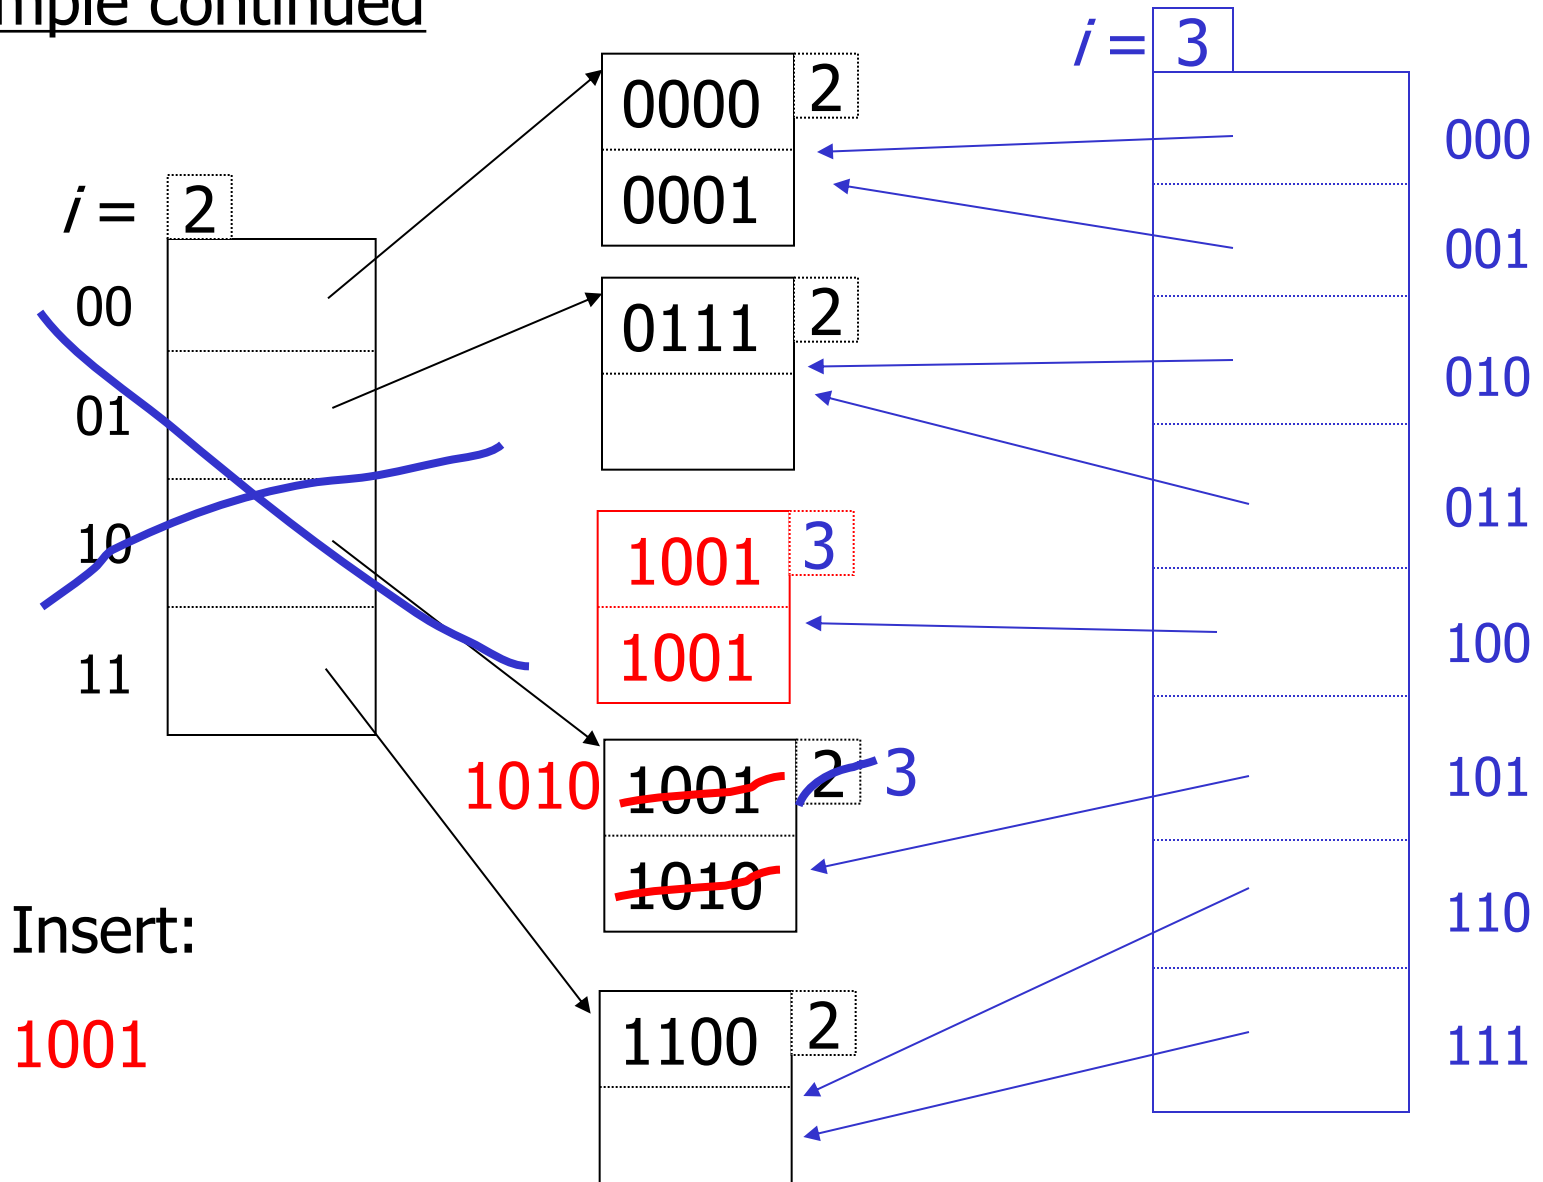
Example continued



Example continued



Example continued



Extensible hashing: deletion

- Merge blocks and cut directory if possible
(Reverse insert procedure)
- Two blocks can be merged after a deletion if all records in these blocks can fit in a single block
- One of these two blocks can be removed and the local depth of the other block one can be decreased by 1
- If **all** local depths are strictly smaller than the global depth i , then the directory can be cut in half and the

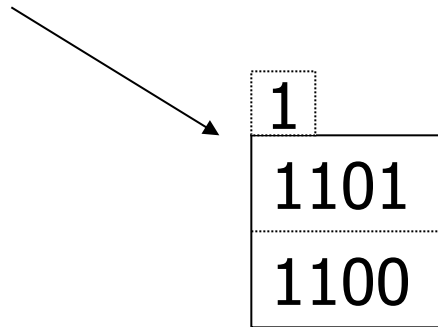
Deletion example:

- Run through insert example in reverse!

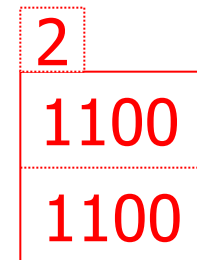
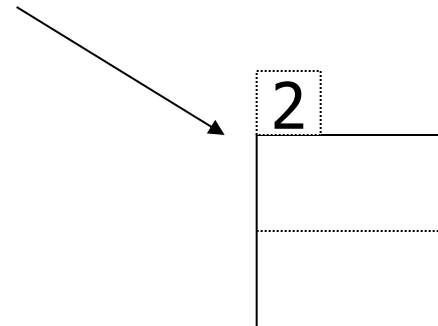
Note: Still need overflow chains

- Example: many records with duplicate keys

insert 1100

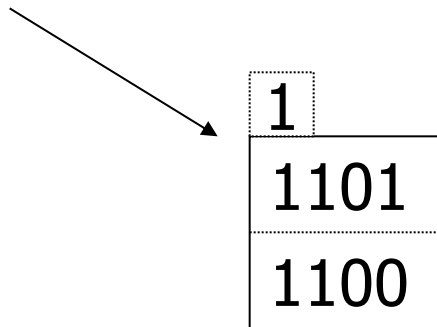


if we split:



Solution: overflow chains

insert 1100



add overflow block:

