

Key-value stores

**The MapReduce Data and Programming Model
(Simulated in Object-Relational SQL)**

**An Introduction to Distributed/Parallel Query Processing Based on
Data Partitioning**

Presented by Muazzam Siddiqui

Slides mainly adapted from Dirk Van Gucht

Key-values stores and queries

- A **key-value store** \mathbf{S} is a binary relation with **schema**
(key: K , value: V)
where K and V are **types** with **domains** $\text{dom}(K)$ and $\text{dom}(V)$ of **objects**
- Note that the key attribute K is not necessarily a primary key of \mathbf{S} . It is possible to have different key-value pairs in \mathbf{S} with the same key value
- A **key-value query** $q : \mathbf{S}_1 \rightarrow \mathbf{S}_2$ is a mapping that sends a key-value store $\mathbf{S}_1(K_1, V_1)$ to a key-value store $\mathbf{S}_2(K_2, V_2)$

Key-value query languages and data-compute environments

- A key-value query q can be programmed in any number of **(database) programming language** such as (Object-Relational) SQL, PhP, MapReduce, Hive, PigLatin, Scala, Python, Java, Javascript, R, scripting languages, etc
- The database programming language is implemented in a **data-compute environment (system)** which can centralized or distributed and which has a certain model of **data storage, data transfer, and communication**: MySQL, PostgreSQL, Hadoop, Spark, MongoDB, Cloudera Impala, Neo4J, Amazon AWS, etc

Key-value query languages and data-compute environments

- Database applications are often implemented and run in different such systems
- Performance can be greatly affected by this choice
- The **database-information ecosystem** is bewildering complex and varied, and keeps changing

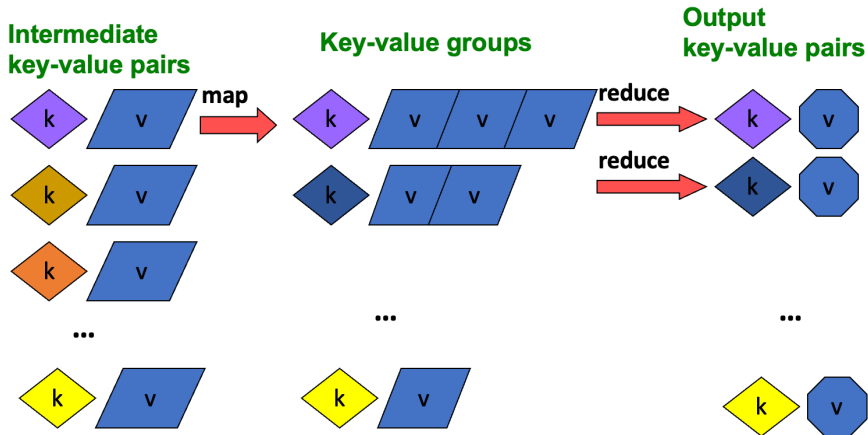
The MapReduce Data and Programming Model

- In this lecture, we will introduce the **MapReduce programming language** as an early example of a **noSQL** language to program key-value pair queries
- MapReduce exhibits a variety of database modeling and query language concepts that have impacted newer database models and query languages

The MapReduce Data and Programming Model

- MapReduce as a programming environment is disappearing but is still very useful to discuss important concepts such as [key-value stores](#), [data partitioning](#), and [distributed query processing](#)
- We will cover the semantics of MapReduce and use the object-relational database model (such as implemented in PostgreSQL) to simulate it

The MapReduce Data and Programming Model



The **Word Count** query (Running Example)

- A key-value store **Document**(doc: text, words: text[]) stores a set of documents
- Each document is represented as a (doc, words) pair:
 - 1 doc is the document identifier of the document
 - 2 words is the set (bag) of words contained in the document

The **Word Count** query (Running Example)

- We wish to implement the **word count** query that maps the key-value store **Document** to the key-value store *wordCount*(*word* : *text*, *wordcount* : *int*)
- The **input** to the **word count** query is the **Document** store
- The **output** of the **word count** query is the store of (*word*, *wordCount*) **pairs** where *wordCount* is the number of occurrences of the word *word* across all the documents in the **Document** store
- In summary, we want to implement

WordCount : *Document*(*text*, *words*) \rightarrow *wordCount*(*word*, *wordcount*)

The **Word Count** query (Input)

- The store **Document** is created as follows:¹

```
CREATE TABLE Document(doc text, words text[]);
```

- Assume that we populate the **Document** store as follows:²

```
INSERT INTO Document VALUES ('d1', ARRAY['A','B','C']),  
                              ('d2', ARRAY['B','C','D']),  
                              ('d3', ARRAY['A','E']),  
                              ('d4', ARRAY['B','B','A','D']),  
                              ('d5', ARRAY['E','F']),
```

Document

doc	words
d1	{A,B,C}
d2	{B,C,D}
d3	{A,E}
d4	{B,B,A,D}
d5	{E,F}

¹Notice that we represent a bag of words with an array.

²Notice that a word may occur multiple times in a document.

The **Word Count** query (Output)

Document

doc	words
d1	{A,B,C}
d2	{B,C,D}
d3	{A,E}
d4	{B,B,A,D}
d5	{E,F}

word count query
→

wordCount

word	wordcount
A	3
B	4
C	2
D	2
E	2
F	1

The **Word Count** query in object-relational SQL

- First formulation:

```
SELECT      p.word, CARDINALITY(ARRAY_AGG(p.doc)) AS wordCount
FROM        (SELECT d.doc, UNNEST(d.words) AS word
              FROM    Document d) p
GROUP BY    (p.word)
```

- Second formulation: Notice that we don't need the actual values of the document identifiers "doc". Thus, we can also formulate the **word count** query as follows:

```
SELECT      p.word, CARDINALITY(ARRAY_AGG(p.one)) AS wordCount
FROM        (SELECT UNNEST(d.words) AS word, 1 as one
              FROM    Document d) p
GROUP BY    (p.word)
```

The **Word Count** query in object-relational SQL (MapReduce style)

- Before we present the MapReduce simulation, we will decompose its object-relational SQL formulation into 3 phases that aligns with the MapReduce programming model
- This formulation will serve as a blueprint for the MapReduce simulation

WITH

%mapper phase:

```
doc_word AS (SELECT UNNEST(d.words) as word, 1 AS one
              FROM Document d),
```

%group (shuffle) phase:

```
word_ones AS (SELECT p.word, ARRAY_AGG(p.one) AS ones
              FROM doc_word p
              GROUP BY (p.word)),
```

%reducer phase:

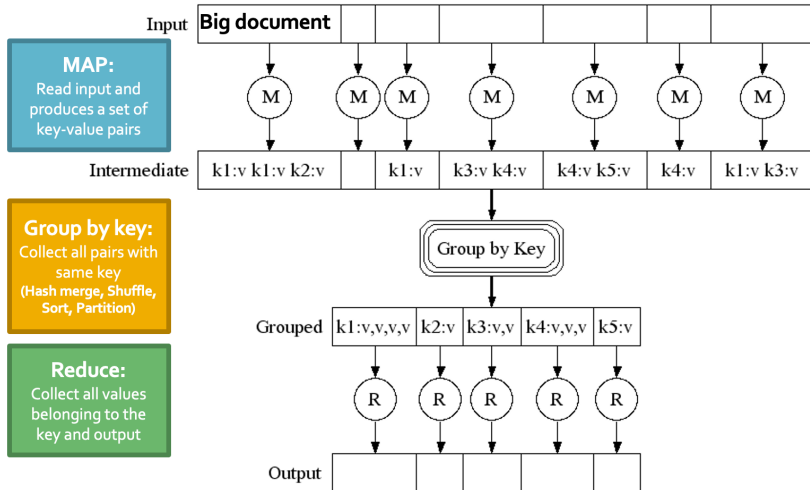
```
word_count AS (SELECT q.word, CARDINALITY(q.ones) AS wordCount
                FROM word_ones q)
```

%output:

SELECT

```
word, wordCount FROM word_count
```

The MapReduce Data and Programming Model



The **Word Count** query (mapper phase)

%mapper phase:

doc_word AS (SELECT **UNNEST**(d.words) as word, 1 AS one
FROM Document d),

Document

doc	words
d1	{A,B,C}
d2	{B,C,D}
d3	{A,E}
d4	{B,B,A,D}
d5	{E,F}

mapper phase
→

doc_word

word	one
A	1
B	1
C	1
B	1
C	1
D	1
A	1
E	1
B	1
B	1
A	1
D	1
E	1
F	1

The **Word Count** query (group (shuffle) phase)

%group (shuffle) phase:

```
word_ones AS (SELECT p.word, ARRAY_AGG(p.one) AS ones
FROM doc_word p
GROUP BY (p.word)),
```

doc_word

word	one
A	1
B	1
C	1
B	1
C	1
D	1
A	1
E	1
B	1
B	1
A	1
D	1
E	1
F	1

group (shuffle) phase
→
group by word and aggregate the 1's

word_ones

word	ones
A	{1,1,1}
B	{1,1,1,1}
C	{1,1}
D	{1,1}
E	{1,1}
F	{1}

The **Word Count** query (reducer phase)

%reducer phase:

```
word_count AS (SELECT q.word, CARDINALITY(q.ones) AS wordCount
FROM word_ones q)
```

word_ones	
word	ones
A	{1,1,1}
B	{1,1,1,1}
C	{1,1}
D	{1,1}
E	{1,1}
F	{1}

reducer phase
→
per word, count its document occurrences

word_count	
word	wordCount
A	3
B	4
C	2
D	2
E	2
F	1

MapReduce queries (mapper and reducer functions)

- A **basic MapReduce query** is a pair of functions:
(**mapper**, **reducer**)
- The **mapper** function takes as input a (*key*, *value*)-pair and outputs a relation (bag) of (*key*, *value*)-pairs:

function **mapper** (key T_1 , value T_2)
returns table(key T_3 , value T_4)

- The **reducer** function takes as input a (*key*, *bagOfValues*)-pair and outputs a relation (bag) of (*key*, *value*)-pairs:

function **reducer** (key T_3 , values $T_4[]$)
returns table(key T_5 , value T_6)

- Notice how the output types T_3 and T_4 of the **mapper** function line up with the input types T_3 and $T_4[]$ of the **reducer** function

MapReduce queries

- A **MapReduce query** is a **sequence** (composition) of basic MapReduce queries

$(mapper_1, reducer_1); (mapper_2, reducer_2); \dots; (mapper_n, reducer_n)$

- The output type of the key-value pairs produced by $reducer_i$ must be the same as the input type for $mapper_{i+1}$, for $i \in [1, n - 1]$

Semantics of a MapReduce query

- The semantics of a **basic** MapReduce query consists of a **mapper-**, a **group (shuffle-)**, and a **reducer-phase**:

- In the **mapper phase**, the mapper is **map-applied**³ to an input relation of (key,value) pairs and the result is put into an intermediate relation *output_map* also consisting of (key,value) pairs
- In the **group (shuffle) phase**, the *output_map* relation is grouped⁴ on its key attribute⁵, and for each key *k*, a pair (*k*, *bagOfValues(k)*) is produced, where

$$\text{bagOfValues}(k) = \{ v \mid \text{output_map}(k, v) \},$$

i.e, the bag of all values in *output_map* with key value *k*

- In the **reducer phase**, the reducer is **map-applied**⁶ to the (*key*, *bagOfValues*) pairs produced in the group (shuffle) phase. For each such pair, the reducer produces a relation of (*key*, *value*) pairs.
- The **output** of the query is the union of all the relations produced in the **reduce_phase**

- The semantics of a MapReduce query is the **composition** of its **basic MapReduce queries**

³Typically in a parallelized and/or distributed manner.

⁴Sometimes also called *shuffled*

⁵Typically by hashing on this key-value.

⁶Typically in a parallelized and/or distributed manner.

Simulating a (basic) MapReduce query in SQL

- At the beginning of the lecture, we formulated the **word count** query in object-relational SQL.
- We deliberately wrote it in a fashion that resembles the mapper-, group (shuffle)-, and reducer phases present in the semantics of a basic MapReduce query
- We can give an even more faithful simulation if we write this SQL query using a **mapper** function and a **reducer** function
- These functions can be programmed as Object-Relational SQL user-defined functions

The **mapper** function

- We formulate the **mapper** function as follows:

```
CREATE FUNCTION mapper(doc text, words text[])  
  RETURNS TABLE (word text, one int) AS  
$$  
  SELECT UNNEST(words) as word, 1 as one;  
$$ LANGUAGE SQL;
```

- The **mapper** function does the following when applied to a document (doc,words) pair:

```
SELECT word, one  
FROM mapper('d1',ARRAY['A','A','B']);
```

word	one
A	1
A	1
B	1

The mapper function (map-apply simulation)

- The **mapper** function, when map-applied to the **Document** relation, produces the relation `map_output(word, one)`.

```
WITH map_output AS
  (SELECT q.word, q.one
   FROM   Document d,
          LATERAL(SELECT p.word, p.one
                   FROM   mapper(d.doc,d.words) p) q)
```

- Notice how we use the **LATERAL** clause
- This is convenient since we need to map-apply the **mapper** function to each document **d** in the **Document** store

The **mapper** function (map-apply)

Document

doc	words
d1	{A,B,C}
d2	{B,C,D}
d3	{A,E}
d4	{B,B,A,D}
d5	{E,F}

map-apply the **mapper** function
→

map_output

word	one
A	1
B	1
C	1
B	1
C	1
D	1
A	1
E	1
B	1
B	1
A	1
D	1
E	1
F	1

Simulation of group (shuffle) phase

- Before we specify the **reducer** function, we show how the **group-phase** is simulated
- This will be done by taking the `map_output` relation and grouping it on the key attribute `word`

```
WITH group_output AS  
  (SELECT p.word, ARRAY_AGG(p.one) as ones  
   FROM   map_output p  
   GROUP BY (p.word))
```

Simulation of the group (shuffle) phase

map_output

word	one
A	1
B	1
C	1
B	1
C	1
D	1
A	1
E	1
B	1
B	1
A	1
D	1
E	1
F	1

group (shuffle)
→

group_output

word	ones
A	{1,1,1}
B	{1,1,1,1}
C	{1,1}
D	{1,1}
E	{1,1}
F	{1}

The **reducer** function

- We now formulate the **reducer** function.
- In our case, this function takes as input a
(word, {1, ..., 1})
pair and outputs the desired
(word, COUNT({1, ..., 1}))
pair (in a relation)

```
CREATE FUNCTION reducer(word text, ones int[])  
  RETURNS TABLE(word text, wordCount int) AS  
$$  
  SELECT word, CARDINALITY(ones) as wordCount;  
$$ LANGUAGE SQL;
```

- The **reducer** function does the following when applied to a
(word,ones) pair:

```
SELECT word, wordCount  
FROM   reducer('A',{1,1,1,1});
```

word	wordCount
A	4

The **reducer** function (map-apply simulation)

- We can now show the simulation of the **reducer phase**:
- I.e, map-apply the **reducer** function to the (word,ones) pairs generated in the **group (shuffle) phase** and produce the output of the word count MapReduce query

```
SELECT r.word, r.wordCount
FROM group_output q,
     LATERAL(SELECT p.word, p.wordCount
              FROM reducer(q.word, r.ones) p) r
```

- So we have

group_output	
word	ones
A	{1,1,1}
B	{1,1,1,1}
C	{1,1}
D	{1,1}
E	{1,1}
F	{1}

map-apply the **reducer** function
→

word	wordCount
A	3
B	4
C	2
D	2
E	2
F	1

The **word count** MapReduce query in Object-Relational SQL

```
CREATE FUNCTION mapper(doc text, words text[])  
  RETURNS TABLE (word text, one int) AS  
$$  
  SELECT UNNEST(words), 1  
$$ LANGUAGE SQL;
```

```
CREATE FUNCTION reducer(word text, ones int[])  
  RETURNS TABLE (word text, wordCount int) AS  
$$  
  SELECT word, CARDINALITY(ones);  
$$ LANGUAGE SQL;
```

MapReduce simulation in Object-Relational SQL

- Putting everything together, we get the following simulation of the **word count** MapReduce query in Object-Relational SQL:

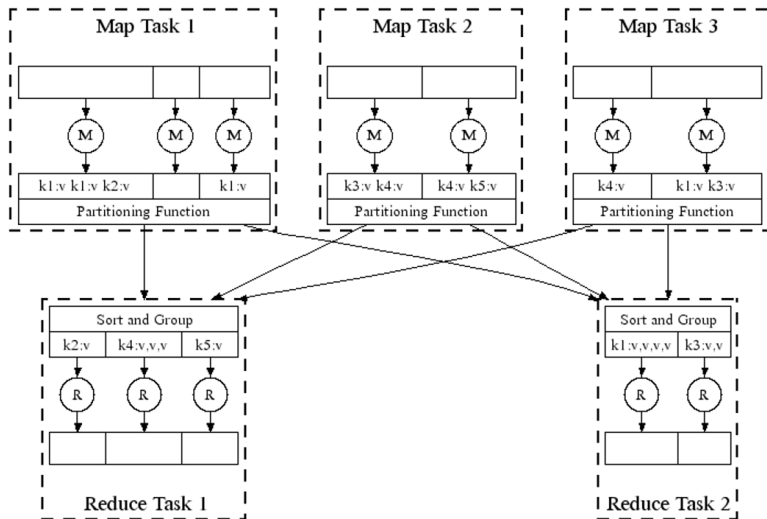
```
WITH
%mapper phase
map_output AS
(SELECT q.word, q.one
 FROM   Document d,
        LATERAL(SELECT p.word, p.one
                  FROM   mapper(d.doc,d.words) p) q),

%group phase
group_output AS
(SELECT p.word, array_agg(p.one) as ones
 FROM   map_output p
 GROUP BY (p.word)),

%reducer phase
reduce_output AS
(SELECT r.word, r.wordCount
 FROM   group_output q,
        LATERAL(SELECT p.word, p.wordCount
                  FROM   reducer(q.word, q.ones) p) r),

%output
SELECT word, wordCount
FROM reduce_output;
```

MapReduce in distributed setting



MapReduce in distributed setting

- In a distributed setting of compute nodes connected by a network, the **Document** key-value store is stored or can be **partitioned** into **chunks** that reside at or can be distributed to the local file systems of the compute nodes
- There are **multiple copies** of the same **mapper** function at the compute nodes these can be evaluated (i.e., map-applied) independently and in parallel (there is no shared memory nor explicit communication)
- More specifically, an instance of the mapper function processes the chunk of **Document** at its compute node and provide its output to the group (shuffle) process
- The group (shuffle) phase is typically implemented by applying a **hash-function** to the keys of the (key,value) pairs emitted by the mappers. Applied to a key, this hash-function will give the location of another compute node
- A (key,value) pair is sent to the compute node with that key's hash-function value. Notice that, because of the properties of hashing functions, **all** (key,value) pairs with the same key will be sent to the same compute node

MapReduce in distributed setting

- There are multiple copies of the **reducer** function waiting at the compute nodes where the hashed (key,value) pairs arrive
- After **all** the appropriate values for a key have been sent and received by the appropriate compute nodes, the reducers can go to work locally (at the compute node) on the list of values associated with a key
- The reducers can transmit their output, or they can keep it locally for further processing by other MapReduce queries.
- A big problem is skew in the data. It is possible that there is an uneven distribution of the values associated with keys. In that case, the computation can be slowed at the reducer compute nodes and the benefits of parallel (distributed) computing can be lost.

The **word count** MapReduce query in Object-Relational SQL (Chunks/Partitioned Data)

- Assume that the document database **Document** is partitioned into a collection of pairwise disjoint document databases **Document**₁, ..., **Document**_n:

$$\mathbf{Document} = \mathbf{Document}_1 \cup \cdots \cup \mathbf{Document}_n.$$

- We can use the same **mapper** function and apply it independently and in parallel to each of the **Document**_i
- The results of these *n* **mapper** computations can be accumulated and then grouped (shuffled) and sent to the **reducer** function.
- Observe that the **mapper** and **reducer** functions do not need to be changed.
- The simulation of this process is shown on the next slide

The **word count** MapReduce query in Object-Relational SQL (Alternative; Partitioned Data)

WITH

%mapper phase— mapper function is evaluated on separate document chunks and accumulated with **UNION**

```
map_output AS
  (SELECT q.word, q.one
   FROM   Document1 d1,
         LATERAL(SELECT p.word, p.one
                  FROM   mapper(d1.doc,d1.words) p)

   UNION
   ...
   UNION
  SELECT q.word, q.one
   FROM   Documentn dn,
         LATERAL(SELECT p.word, p.one
                  FROM   mapper(dn.doc,dn.words) p) q)
```

%group phase

```
group_output AS
  (SELECT p.word, array_agg(p.one) as ones
   FROM   map_output p
   GROUP BY (p.word)),
```

%reducer phase

```
reduce_output AS
  (SELECT r.word, r.wordCount
   FROM   group_output q,
         LATERAL(SELECT p.word, p.wordCount
                  FROM   reducer(q.word, q.ones) p) r)
```

%output

```
SELECT word, wordCount
FROM reduce_output;
```

MapReduce Limitations

- **Limitation 1:**
 - Notice that a MapReduce query does not have conditional statements nor loop statements
 - This is a serious limitation for problems that require iteration to find solutions
 - Such is the case in many data science and machine learning problems (for example page-rank algorithm, k-means clustering, gradient-descent in deep learning etc.)
- **Limitation 2:** Performance can be greatly affected by **skew in data**: during the grouping, different keys may be associated with vastly different numbers of values

MapReduce Limitations

■ Limitation 3:

- Between successive basic MapReduce query, intermediate results are written to files
- This limits in-memory processing

■ Limitation 4:

- The input to a MapReduce program is a **single** relation.
- Notice that therefore to do a **binary** operation such a join, union, intersection, or set-difference between two relations R and S , it is necessary store/model the data in both R and S into a **single** relation of key-values pairs— **awkward data modeling**