


Concurrency Control

Based on notes by Hector Garcia-
Molina

Concurrency Control



Transactions execute
concurrently

Each transaction has own
local memory

Transactions do not
Communicate

Example:

T1: Read(A)
A \leftarrow A+100
Write(A)
Read(B)
B \leftarrow B+100
Write(B)

T2: Read(A)
A \leftarrow A \times 2
Write(A)
Read(B)
B \leftarrow B \times 2
Write(B)

Constraint: A=B

Schedule A

T1

T2

Read(A); $A \leftarrow A + 100$

Write(A);

Read(B); $B \leftarrow B + 100$;

Write(B);

Read(A); $A \leftarrow A \times 2$;

Write(A);

Read(B); $B \leftarrow B \times 2$;

Write(B);

Schedule A

T1	T2	A	B
		25	25
Read(A); $A \leftarrow A + 100$			
Write(A);		125	
Read(B); $B \leftarrow B + 100$;			
Write(B);			125
	Read(A); $A \leftarrow A \times 2$;		
	Write(A);	250	
	Read(B); $B \leftarrow B \times 2$;		
	Write(B);		250
		250	250

Schedule B

T1

T2

Read(A); $A \leftarrow A \times 2$;

Write(A);

Read(B); $B \leftarrow B \times 2$;

Write(B);

Read(A); $A \leftarrow A + 100$

Write(A);

Read(B); $B \leftarrow B + 100$;

Write(B);

Schedule B

T1	T2	A	B
		25	25
	Read(A); $A \leftarrow A \times 2$;		
	Write(A);	50	
	Read(B); $B \leftarrow B \times 2$;		50
	Write(B);		
Read(A); $A \leftarrow A + 100$			
Write(A);		150	
Read(B); $B \leftarrow B + 100$;			150
Write(B);			
		150	150

Schedule C

T1

T2

Read(A); $A \leftarrow A + 100$

Write(A);

Read(A); $A \leftarrow A \times 2$;

Write(A);

Read(B); $B \leftarrow B + 100$;

Write(B);

Read(B); $B \leftarrow B \times 2$;

Write(B);

Schedule C

T1	T2	A	B
		25	25
Read(A); $A \leftarrow A + 100$			
Write(A);		125	
	Read(A); $A \leftarrow A \times 2$;		
	Write(A);	250	
Read(B); $B \leftarrow B + 100$;			125
Write(B);			
	Read(B); $B \leftarrow B \times 2$;		
	Write(B);		250
		250	250

Schedule D

T1

T2

Read(A); $A \leftarrow A + 100$

Write(A);

Read(A); $A \leftarrow A \times 2$;

Write(A);

Read(B); $B \leftarrow B \times 2$;

Write(B);

Read(B); $B \leftarrow B + 100$;

Write(B);

Schedule D

T1	T2	A	B
		25	25
Read(A); $A \leftarrow A + 100$			
Write(A);		125	
	Read(A); $A \leftarrow A \times 2$;		
	Write(A);	250	
	Read(B); $B \leftarrow B \times 2$;		
	Write(B);		50
Read(B); $B \leftarrow B + 100$;			
Write(B);			150
		250	150

Schedule E

Same as Schedule D
but with new T2'

T1

T2'

Read(A); $A \leftarrow A + 100$

Write(A);

Read(A); $A \leftarrow A \times 1$;

Write(A);

Read(B); $B \leftarrow B \times 1$;

Write(B);

Read(B); $B \leftarrow B + 100$;

Write(B);

Schedule E

Same as Schedule D
but with new T2'

T1	T2'	A	B
		25	25
Read(A); $A \leftarrow A + 100$			
Write(A);		125	
	Read(A); $A \leftarrow A \times 1$;		
	Write(A);	125	
	Read(B); $B \leftarrow B \times 1$;		
	Write(B);		25
Read(B); $B \leftarrow B + 100$;			
Write(B);			125
		125	125


- Want schedules that are “good”, regardless of
 - initial state and
 - transaction semantics
- Only look at order of read and writes

Example:

$Sc = r_1(A)w_1(A)r_2(A)w_2(A)r_1(B)w_1(B)r_2(B)w_2(B)$

Example:

$$Sc = r_1(A)w_1(A)r_2(A)w_2(A)r_1(B)w_1(B)r_2(B)w_2(B)$$



$$Sc' = r_1(A)w_1(A) \ r_1(B)w_1(B)r_2(A)w_2(A)r_2(B)w_2(B)$$

T_1

T_2

However, for S_d :

$$S_d = \underbrace{r_1(A)w_1(A)}_{\text{Group 1}} \underbrace{r_2(A)w_2(A)}_{\text{Group 2}} \underbrace{r_2(B)w_2(B)}_{\text{Group 3}} \underbrace{r_1(B)w_1(B)}_{\text{Group 4}}$$

The diagram illustrates a non-commutative relationship between the groups. Two curved arrows are drawn below the equation. The first arrow starts under the first group, $r_1(A)w_1(A)$, and points to the third group, $r_2(B)w_2(B)$. The second arrow starts under the fourth group, $r_1(B)w_1(B)$, and points to the second group, $r_2(A)w_2(A)$. Both arrows are marked with a large red 'X', indicating that the groups do not commute.

Concepts

Transaction: sequence of $r_i(x)$, $w_i(x)$ actions

Conflicting actions:

	$r_i(A)$	$w_i(A)$	$w_i(A)$	
	\swarrow	\swarrow	\swarrow	
	$w_j(A)$	$r_j(A)$	$w_j(A)$	$i \neq j$

Schedule: represents chronological order
in which actions are executed

Serial schedule: no interleaving of actions
or transactions

Definition

S_1, S_2 are conflict equivalent schedules if S_1 can be transformed into S_2 by a series of swaps on non-conflicting actions.

Definition


S_1 is serializable if there is a serial schedule S_s such that for every initial database state, the effects of S_1 and S_s are the same

Definition

A schedule is conflict serializable if it is conflict equivalent to some serial schedule.

Example: **Sc is conflict serializable**

$Sc = r_1(A)w_1(A)r_2(A)w_2(A)r_1(B)w_1(B)r_2(B)w_2(B)$



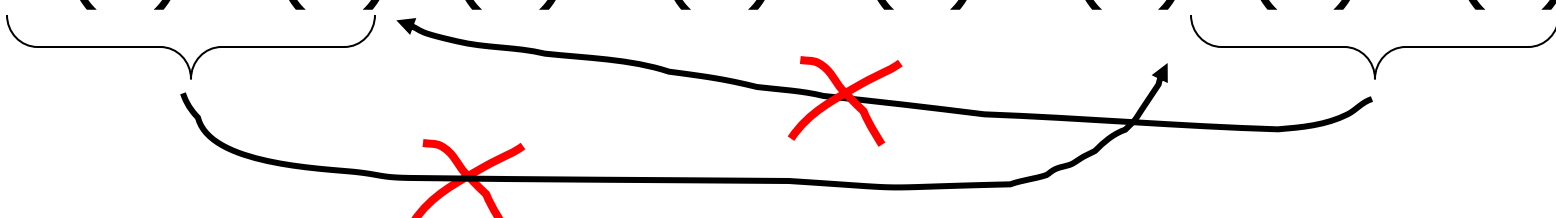
$Sc' = r_1(A)w_1(A) r_1(B)w_1(B)r_2(A)w_2(A)r_2(B)w_2(B)$

T_1

T_2

Sd is **not conflict serializable**:

$Sd = r_1(A)w_1(A)r_2(A)w_2(A) r_2(B)w_2(B)r_1(B)w_1(B)$



Precedence graph $P(S)$ (S is schedule)


Nodes: transactions in S

Arcs: $T_i \rightarrow T_j$ whenever

- $p_i(A), q_j(A)$ are actions in S
- $p_i(A) <_S q_j(A)$
- at least one of p_i, q_j is a write

Exercise:


- What is $P(S)$ for
 $S = w_3(A) w_2(C) r_1(A) w_1(B) r_1(C) w_2(A) r_4(A) w_4(D)$



- Is S serializable?

Another Exercise:

- What is $P(S)$ for
 $S = w_1(A) \ r_2(A) \ r_3(A) \ w_4(A)$?



Lemma

S_1, S_2 conflict equivalent $\Rightarrow P(S_1)=P(S_2)$

Lemma

S_1, S_2 conflict equivalent $\Rightarrow P(S_1)=P(S_2)$

Proof:

Assume $P(S_1) \neq P(S_2)$

$\Rightarrow \exists T_i: T_i \rightarrow T_j$ in S_1 and not in S_2

$\Rightarrow S_1 = \dots p_i(A) \dots q_j(A) \dots$	$\left\{ \begin{array}{l} p_i, q_j \\ \text{conflict} \end{array} \right.$
$S_2 = \dots q_j(A) \dots p_i(A) \dots$	

$\Rightarrow S_1, S_2$ not conflict equivalent

Note: $P(S_1)=P(S_2) \not\Rightarrow S_1, S_2$ conflict equivalent

Note: $P(S_1)=P(S_2) \not\Rightarrow S_1, S_2$ conflict equivalent

Counter example:

$T1 := w1(A) \ r1(B)$

$T2 := r2(A) \ w2(B)$

$S_1 = w1(A) \ r2(A) \ w2(B) \ r1(B)$

$S_2 = r2(A) \ w1(A) \ r1(B) \ w2(B)$



Serialization Theorem

$P(S_1)$ acyclic $\iff S_1$ conflict serializable

Serialization Theorem

$P(S_1)$ acyclic $\iff S_1$ conflict serializable

(\Leftarrow) Assume S_1 is conflict serializable


$\Rightarrow \exists S_s: S_s, S_1$ conflict equivalent

$\Rightarrow P(S_s) = P(S_1)$

$\Rightarrow P(S_1)$ acyclic since $P(S_s)$ is acyclic

Serialization Theorem

$P(S_1)$ acyclic $\iff S_1$ conflict serializable



Serialization Theorem

$P(S_1)$ acyclic $\iff S_1$ conflict serializable


(\implies) Assume $P(S_1)$ is acyclic

Transform S_1 as follows:

(1) Take T_1 to be transaction with no incident arcs


(2) Move all T_1 actions to the front

$S_1 = \dots q_j(A) \dots p_1(A) \dots$



(3) we now have $S_1 = \langle T_1 \text{ actions} \rangle \langle \dots \text{rest} \dots \rangle$

(4) repeat above steps to serialize rest!



Serialization Theorem


$P(S)$ acyclic $\Leftrightarrow S$ conflict serializable

How to enforce serializable schedules?

Option 1: run system, recording $P(S)$;
after some time,
check for $P(S)$ cycles and declare
if execution was good

How to enforce serializable schedules?

Option 2: prevent P(S) cycles from occurring




A locking protocol

Two new actions:

lock (exclusive): $li(A)$

unlock: $ui(A)$



Rule #1: Well-formed transactions

$T_i: \dots l_i(A) \dots p_i(A) \dots u_i(A) \dots$

A transaction can only read or write if it first as been granted a lock.

If a transaction locks an element, it must later unlock it.

Rule #2 Legal scheduler

$$S = \dots\dots li(A) \dots\dots ui(A) \dots\dots$$

\longleftrightarrow
no $lj(A)$

Locks must have their intended semantics: no 2 transactions may have locked the same element without one having first released the lock.

Exercise:

- What schedules are legal?

What transactions are well-formed?

$S1 = l_1(A)l_1(B)r_1(A)w_1(B)l_2(B)u_1(A)u_1(B)$
 $r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

$S2 = l_1(A)r_1(A)w_1(B)u_1(A)u_1(B)$
 $l_2(B)r_2(B)w_2(B)l_3(B)r_3(B)u_3(B)$

$S3 = l_1(A)r_1(A)u_1(A)l_1(B)w_1(B)u_1(B)$
 $l_2(B)r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

Exercise:

- What schedules are legal?

What transactions are well-formed?

$S1 = l_1(A)l_1(B)r_1(A)w_1(B)l_2(B)u_1(A)u_1(B)$
 $r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

$S2 = l_1(A)r_1(A)w_1(B)u_1(A)u_1(B)$
 $l_2(B)r_2(B)w_2(B)l_3(B)r_3(B)u_3(B)u_2(B)?$

$S3 = l_1(A)r_1(A)u_1(A)l_1(B)w_1(B)u_1(B)$
 $l_2(B)r_2(B)w_2(B)u_2(B)l_3(B)r_3(B)u_3(B)$

Do rules 1 and 2 guarantee
serializability?

Schedule F

T1

$l_1(A); \text{Read}(A)$

$A \leftarrow A + 100; \text{Write}(A); u_1(A)$

$l_1(B); \text{Read}(B)$

$B \leftarrow B + 100; \text{Write}(B); u_1(B)$

T2

$l_2(A); \text{Read}(A)$

$A \leftarrow A \times 2; \text{Write}(A); u_2(A)$

$l_2(B); \text{Read}(B)$


$B \leftarrow B \times 2; \text{Write}(B); u_2(B)$

Schedule F


		A	B
T1	T2	25	25
l ₁ (A);Read(A)			
A ← A + 100;Write(A);u ₁ (A)		125	
	l ₂ (A);Read(A)		
	A ← A x 2;Write(A);u ₂ (A)	250	
	l ₂ (B);Read(B)		
	B ← B x 2;Write(B);u ₂ (B)		50
l ₁ (B);Read(B)			
B ← B + 100;Write(B);u ₁ (B)			150
		250	150

Rule #3 Two phase locking (2PL) for transactions


$T_i = \dots \dots \text{li}(A) \dots \dots \text{ui}(A) \dots \dots$



no unlocks



no locks



Schedule G


T1	T2
$l_1(A); \text{Read}(A)$	
$A \leftarrow A + 100; \text{Write}(A)$	
$l_1(B); u_1(A)$	
	$l_2(A); \text{Read}(A)$
	$A \leftarrow A \times 2; \text{Write}(A); l_2(B)$

delayed

Schedule G

T1	T2
$l_1(A); \text{Read}(A)$	
$A \leftarrow A + 100; \text{Write}(A)$	
$l_1(B); u_1(A)$	
	$l_2(A); \text{Read}(A)$
	$A \leftarrow A \times 2; \text{Write}(A); l_2(B)$
$\text{Read}(B); B \leftarrow B + 100$	
$\text{Write}(B); u_1(B)$	


delayed



Schedule G

T1	T2
$l_1(A); \text{Read}(A)$	
$A \leftarrow A + 100; \text{Write}(A)$	
$l_1(B); u_1(A)$	
	$l_2(A); \text{Read}(A)$ delayed
	$A \leftarrow A \times 2; \text{Write}(A); l_2(B)$
$\text{Read}(B); B \leftarrow B + 100$	
$\text{Write}(B); u_1(B)$	
	$l_2(B); u_2(A); \text{Read}(B)$
	$B \leftarrow B \times 2; \text{Write}(B); u_2(B);$

Schedule H (T₂ reversed)



Schedule is deadlocked since neither T1 nor T2 can proceed!

2PL guarantees serializability

Show that rules #1,2,3 \Rightarrow conflict-serializable schedules

Conflict rules for $l_i(A)$, $u_i(A)$:

- $l_i(A)$, $l_j(A)$ conflict
- $l_i(A)$, $u_j(A)$ conflict

Note: no conflict $\langle u_i(A), u_j(A) \rangle$, $\langle l_i(A), r_j(A) \rangle$, ...

Theorem Rules #1,2,3 \Rightarrow conflict
(2PL) serializable
schedule

Theorem Rules #1,2,3 \Rightarrow conflict
(2PL) serializable
schedule

To help in proof:

Definition Shrink(T_i) = SH(T_i) =
first unlock action of T_i

Lemma

$$Ti \rightarrow Tj \text{ in } S \Rightarrow SH(Ti) <_S SH(Tj)$$

Lemma

$$T_i \rightarrow T_j \text{ in } S \Rightarrow SH(T_i) <_S SH(T_j)$$

Proof of lemma:

$T_i \rightarrow T_j$ means that

$$S = \dots p_i(A) \dots q_j(A) \dots; \quad p, q \text{ conflict}$$

By rules 1,2:

$$S = \dots p_i(A) \dots u_i(A) \dots l_j(A) \dots q_j(A) \dots$$

Lemma


$$T_i \rightarrow T_j \text{ in } S \Rightarrow SH(T_i) <_S SH(T_j)$$

Proof of lemma:

$T_i \rightarrow T_j$ means that

$$S = \dots p_i(A) \dots q_j(A) \dots; \quad p, q \text{ conflict}$$

By rules 1,2:

$$S = \dots p_i(A) \dots u_i(A) \dots l_j(A) \dots q_j(A) \dots$$


$$\text{By rule 3:} \quad SH(T_i) \qquad SH(T_j)$$

$$\text{So, } SH(T_i) <_S SH(T_j)$$

Theorem Rules #1,2,3 \Rightarrow conflict
(2PL) serializable
schedule

Proof:

(1) Assume $P(S)$ has cycle

$$T_1 \rightarrow T_2 \rightarrow \dots T_n \rightarrow T_1$$

(2) By lemma: $SH(T_1) < SH(T_2) < \dots < SH(T_1)$

(3) Impossible, so $P(S)$ acyclic

(4) $\Rightarrow S$ is conflict serializable